

THE GEORGE
WASHINGTON
UNIVERSITY

WASHINGTON, DC

DATA SCIENCE PROGRAM

CAPSTONE REPORT - SPRING 2024

A Benchmark for Graph-Based Dynamic Recommendation Systems

Tyler Wallett

supervised by
Amir Jafari

Abstract

The surge of Graph Neural Networks (GNNs) has catalyzed significant advancements in recommendation systems by enabling more effective modeling of user-item interactions within undirected bipartite graphs. However, the proliferation of GNN architectures, coupled with the absence of standardized benchmarking frameworks, presents challenges in systematically evaluating and comparing different dynamic recommendation models. In response, we propose a comprehensive benchmarking study of bipartite GNN operators for recommendation systems using the PyTorch Geometric library. Our contributions include the development of a flexible benchmarking framework encompassing data preprocessing, model training, and evaluation protocols, facilitating fair comparison across diverse dynamic recommendation scenarios. We rigorously assess the performance of various GNN models, ranging from traditional methods to state-of-the-art architectures, on the MovieLens100k dataset. Through insightful analysis of experimental results, we elucidate the strengths and weaknesses of different GNN operators and offer practical suggestions for model selection and configuration. Our work aims to foster transparency, reproducibility, and innovation in GNN-based dynamic recommendation systems, providing a valuable resource for researchers and practitioners in the field.

Contents

1. Introduction	3
2. Methodology	3
2.1. Graph Theory	3
2.2. PyTorch Geometric	4
2.3. Models	7
2.3.1. SimpleConv	7
2.3.2. SAGEConv[1]	8
2.3.3. GraphConv[2]	8
2.3.4. ResGatedGraphConv[3]	8
2.3.5. GATConv[4]	8
2.3.6. GATv2Conv[5]	9
2.3.7. TransformerConv[6]	9
2.3.8. GINConv[7]	10
2.3.9. EdgeConv[8]	10
2.3.10. FeaStConv[9]	10
2.3.11. LEConv[10]	11
2.3.12. GENConv[11]	11
2.3.13. WLConvContinuous[12]	11
2.4. Dataset	11
3. Results	12
4. Discussion	14
5. Conclusion	14
A. Summary of Models	15
B. Summary of Hyperparameters	17
C. Summary of Results	18

1. Introduction

In the realm of recommendation systems, the efficacy of modern algorithms is often evaluated on their ability to understand complex user-item interactions within undirected bipartite graphs. These interactions, which represent the relationships between users and items, are pivotal for generating accurate recommendations in various domains such as e-commerce, social media, and content streaming platforms. With the advent of Graph Neural Networks (GNNs), there has been a surge of interest in leveraging their expressive power to model and predict user preferences more effectively.

However, the landscape of GNN-based recommendation models is vast, encompassing diverse architectures, training strategies, and evaluation metrics. Amidst this diversity, there is a pressing need for a standardized benchmarking framework that enables fair comparison and evaluation of different GNN operators tailored specifically for undirected bipartite recommendation graphs.

In response to this need, we present a comprehensive benchmarking study of bipartite GNN operators for dynamic recommendation systems. Leveraging the PyTorch Geometric library, we meticulously evaluate a spectrum of GNN models, ranging from traditional methods to state-of-the-art architectures.

Thus, our contributions are threefold:

1. Benchmarking Framework Development: We propose a flexible and extensible benchmarking framework that facilitates the systematic evaluation of GNN-based recommendation models. This framework encompasses data preprocessing, model training, evaluation protocols, and result analysis, thereby providing a standardized basis for fair comparison.

2. Model Selection and Evaluation: We rigorously assess the performance of various bipartite GNN operators, including GraphSAGE, Graph Attention Networks (GATs), and more, under different experimental settings. By conducting extensive experiments on MovieLens100k dataset, we elucidate the strengths and weaknesses of each model in capturing intricate user-item interactions.

3. Insightful Analysis and Recommendations: Through thorough analysis of experimental results, we uncover key insights into the behavior and performance of different bipartite GNN operators in dynamic recommendation scenarios. Based on these insights, we offer practical suggestions for selecting and configuring GNN models tailored to specific application domains.

By providing a comprehensive benchmarking study along with an open-source implementation in PyTorch Geometric, we aim to foster transparency, reproducibility, and innovation in the development of GNN-based dynamic recommendation systems. We believe that our work will serve as a valuable resource for researchers and practitioners seeking to leverage the power of GNNs for building more effective recommendation engines.

2. Methodology

The purpose of this section is to delineate the structured approach utilized for benchmarking PyTorch Geometric’s bipartite GNN operators for dynamic recommendation systems. To fulfill this purpose this section introduces graph theory concepts for edge-level prediction for dynamic recommendation systems, provides an overview of the PyTorch Geometric pipeline, describes each model, and accounts for the dataset and metrics utilized.

2.1. Graph Theory

Graphs, denoted by $G(V, E)$, consist of two finite sets, V and E . Each element of $V = \{v_1, v_2, \dots, v_n\}$ is called a vertex (node) and represents the entities within the graph. Each el-

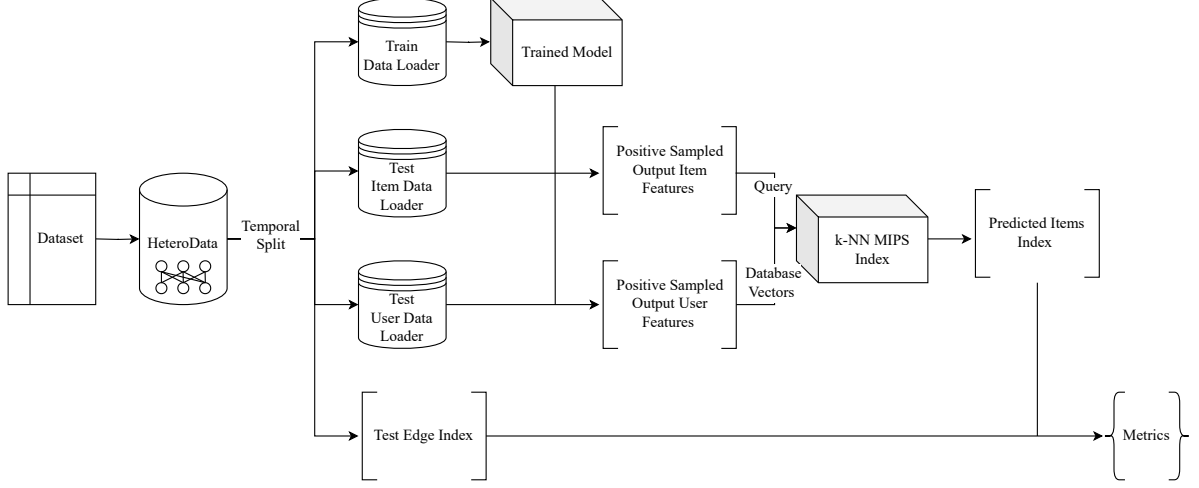


Figure 1: PyTorch Geometric Dynamic Recommendation System Pipeline.

ement of $E = \{\{v_1, v_2\}, \dots, \{v_1, v_n\}\}$ is called an edge (link) and represents the relationships between each vertex within the graph. To fully grasp the graph used for dynamic edge-level prediction for recommendation systems, we first need to understand three types of graphs:

- **Undirected graphs:** Graphs that lack inherent edge directionality, signifying mutual or symmetrical relationships between vertices.
- **Bipartite graphs:** Graphs, denoted as $G(U, V, E)$, that partition vertices into two disjoint sets U and V , ensuring non-adjacency within sets and edges connecting vertices from distinct vertex sets.
- **Heterogeneous graphs:** Graphs that contain either two or more types of vertex sets or two or more types of edge sets.

Hence, for dynamic edge-level prediction for recommendation systems, we adopt an undirected bipartite graph, represented as $G(U, I, E)$, where U denotes users and I denotes recommended items. From this graph, the goal is to predict edges between users and items over time.

2.2. PyTorch Geometric

PyTorch Geometric is a powerful Python library used for developing GNNs that is seamlessly integrated with PyTorch. It offers intuitive functionalities for heterogeneous graph learning and data structures, data loaders, bipartite GNN operators, and k -Nearest Neighbors Maximum Inner Product Search (k -NN MIPS). PyTorch Geometric plays an indispensable role, as the entire benchmark project relies on its pipeline 1. You can find the code used for this benchmark here: <https://github.com/twalleth/DynamicRecSys>.

Now, let's delve into a detailed explanation of each component comprising the PyTorch Geometric pipeline, outlining the structure of the evaluation process:

- **HeteroData:** HeteroData is a data structure in PyTorch Geometric that encompasses each component of a heterogeneous graph. Utilizing this data structure is the first step of the pipeline 1 as we are interested in representing our comma-separated values (csv) dataset as a heterogeneous graph. Specifically, for dynamic edge-level prediction for recommendation systems, we are interested in representing the previously mentioned undirected bipartite graph $G(U, I, E)$.

The structure of $G(U, I, E)$ is characterized by three main components. Each user vertex inside the user vertex set U is represented by an embedding vector, or vertex feature, $U = [\mathbf{x}_{user_1}, \mathbf{x}_{user_2}, \dots, \mathbf{x}_{user_{|U|}}]$. These user embeddings are first initialized with an identity matrix I of dimension $|U| \times |U|$. Similarly, each item vertex inside the item vertex set I is represented by an embedding vector, or vertex feature, $I = [\mathbf{x}_{item_1}, \mathbf{x}_{item_2}, \dots, \mathbf{x}_{item_{|I|}}]$. However, the only difference is that these item embeddings are initialized with normal distributed random numbers in a matrix of size $|I| \times Latent Dimension$, where the latent dimension is a hyperparameter. Lastly, edges are represented by a bidirectional edge index (adjacency matrix) with feature and time labels for user-feature-item and item-feature-user relationships.

- **Data Loaders:** In the pipeline 1, before constructing the data loaders there is a temporal split of the dataset. To do so, we simply extract the HeteroData time labels and sort them to be split by a set threshold.

LinkNeighborhood Loader is the data loader used for training the models. This data loader samples bipartite edges from the graph $G(U, I, E)$ with temporal considerations. It specifies parameters such as the number of neighbors to be sampled, batch size, and temporal attributes in the graph data. Additionally, it sets up edge labels and their associated indices, while controlling the ratio of negative to positive edges to be sampled during training.

Neighborhood Loader[1] is the data loader used for testing the models. This data loader samples unidirectional edges from the graph $G(U, I, E)$ with temporal considerations, and handles the same parameters as previously mentioned for the LinkNeighborhood Loader.

Note that, in the pipeline 1, a test edge index is also created. This test edge index will be compared to the predicted item’s index to prepare the link-prediction metrics.

- **Heterogeneous Graph Learning:** Heterogeneous Graph Learning refers to any GNN learning on heterogeneous graph data. One of the challenges faced is that Standard Message Passing GNNs (MP-GNNs) can not trivially be applied to heterogeneous graph data, as vertex and edge features from different types can not be processed by the same functions due to differences in feature type. Luckily, PyTorch Geometric provides functionality to automatically convert a homogeneous GNN model to a heterogeneous GNN model.

In the pipeline 1 the design of the trained model is standardized for all models. As the LinkNeighborhood Loader samples training edges in a bidirectional and temporal manner, it also samples corresponding input embeddings. For dynamic edge-level prediction for recommendation systems, these sampled input embeddings can be from either the user \mathbf{x}_{user_j} or item \mathbf{x}_{item_j} . Once these input embeddings are sampled they are passed through a bipartite GNN model, or simply viewed as an **encoder** function:

$$\mathbf{x}'_{item_i} = \text{encoder}(\mathbf{x}_{item_j})$$

$$\mathbf{x}'_{user_i} = \text{encoder}(\mathbf{x}_{user_j})$$

After receiving output embeddings for both user \mathbf{x}'_{user_i} and item \mathbf{x}'_{item_i} , these are then passed through a **decoder** function to return an output vector \mathbf{a} (Note that the choice of **decoder** function for this benchmark was simply an inner product between the two sampled output embeddings):

$$\mathbf{a} = \text{decoder}(\mathbf{x}'_{item_i}, \mathbf{x}'_{user_i})$$

Lastly, the output \mathbf{a} of this **decoder** function is used to calculate a mean squared error e relative to the sampled edge labels \mathbf{t} :

$$e = \frac{1}{n} \sum_{i=1}^n (\mathbf{t} - \mathbf{a})^2$$

By following this procedure, that applies heterogeneous graph learning functionality, we are able to standardize the training of all of the bipartite GNN models.

- **k -NN MIPS:** Maximum Inner Product Search (MIPS) is an algorithmic technique designed to efficiently find an optimal database vector \mathbf{x}_* by maximizing the inner product similarity measure between a corresponding database vector \mathbf{x} and query vector \mathbf{q} . There are many variations of MIPS, one of them being Nearest Neighbor MIPS[13, 14] attempting to find the optimal database vector \mathbf{x}_* by minimizing over a distance function ρ :

$$\mathbf{x}_* = \operatorname{argmin} \rho(\mathbf{q}^T \mathbf{x})$$

k -Nearest Neighbor MIPS[15] is simply an extension of Nearest Neighbor MIPS that attempts to find optimal spatial relationships within a neighborhood of k expansions. PyTorch Geometric provides k -Nearest Neighbor MIPS[15] integration by means of the Faiss library[16]. Therefore, we are able to calculate k optimal spatial relationships between users and items at inference within our pipeline 1.

- **Metrics:** In the evaluation of dynamic recommendation systems, various metrics play crucial roles in quantifying the effectiveness and performance of the models. These metrics provide insights into how well the system is able to suggest relevant items to users.

Recall@ k and Precision@ k are fundamental measures that assess the accuracy and completeness of recommendations within the top- k suggestions, respectively.

$$\text{Recall@}k = \frac{\text{Relevant items in top-}k \text{ recommendations}}{\text{Total number of relevant items}}$$

$$\text{Precision@}k = \frac{\text{Relevant items in top-}k \text{ recommendations}}{k}$$

Mean Average Precision@ k (MAP@ k) takes into account the precision of recommendations at different cutoff levels, providing a holistic view of the system’s performance.

$$\text{MAP@}k = \frac{1}{|Q|} \sum_{q=1}^{|Q|} \frac{1}{\min(k, N_q)} \sum_{i=1}^k P(i)$$

F1@ k offers a balanced evaluation by considering both precision and recall at the k -th recommendation.

$$F1@k = 2 \times \frac{\text{Precision@}k \times \text{Recall@}k}{\text{Precision@}k + \text{Recall@}k}$$

Normalized Discounted Cumulative Gain@ k (nDCG@ k) evaluates the ranking quality of recommended items up to the k -th recommendation, considering both relevance and position in the list.

$$\begin{aligned}
\text{DCG}@k &= \sum_{i=1}^k \frac{2^{rel_i} - 1}{\log_2(i + 1)} \\
\text{IDCG}@k &= \sum_{i=1}^{\min(k, \text{num_relevant})} \frac{2^{rel_i} - 1}{\log_2(i + 1)} \\
\text{nDCG}@k &= \frac{\text{DCG}@k}{\text{IDCG}@k}
\end{aligned}$$

Together, these metrics offer a comprehensive assessment of dynamic recommendation systems, crucial for optimizing their performance and enhancing user satisfaction.

It is important to mention that the early stopping callback was configured with a patience parameter of 20, ensuring that training ceased when there was no improvement in performance over 20 consecutive epochs. The validation score was conducted based on Recall@ k since we prioritize the identification of all relevant items, thus emphasizing the system’s ability to minimize the likelihood of overlooking potentially suitable recommendations. This choice aligns with the practical objectives of our research, which emphasize the importance of providing high-quality recommendations to users.

2.3. Models

In total 13 bipartite GNN operators from Pytorch Geometric were trained. Each model was trained for 1000 epochs using a batch size of 1024 and Stochastic Gradient Descent (SGD) optimization with a learning rate of 3e-03. Further, each model operates with a latent dimension of 64 and samples neighbors with a proximity of [2,2] for a equivalent benchmarking. For specific information regarding the hyperparameters used for the training of each of these models refer to the summary table in the appendix section 3.

Several bipartite GNN models were not included in the benchmarking. Models that necessitated edge labels, our targets \mathbf{t} , to be passed as inputs were obviously excluded, these are: GINEConv[17], GMMConv[18], SplineConv[19], NNConv[20], CGConv[21], and GeneralConv[22]. Lastly, models that required building a different pipeline were also excluded, such as MFConv[23] and FANConv[24].

Additionally, for each bipartite GNN operator, random seeds were set and temporal sampling was run on central processing units (CPUs), thus ensuring the reproducibility of results.

2.3.1. SimpleConv

The SimpleConv (Simple Convolutional Layer) model operates by aggregating information from neighboring items or users. It adopts a simple aggregation strategy where the embedding vector of an item or user is updated by aggregating the embeddings of its neighboring items or users using an aggregation operator. In mathematical terms, for each item or user, denoted as i , its updated embedding, \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} , is obtained by aggregating the embeddings of its neighbors, denoted as $j \in \mathcal{N}(i)$, where $\mathcal{N}(i)$ represents the set of neighbors of i , and \oplus denotes the SetTransformerAggregation[25] operation. This aggregation operator was chosen for this benchmark as it yielded the best results out of all of the aggregation operators.

$$\mathbf{x}'_{item_i} = \oplus_{j \in \mathcal{N}(i)} \mathbf{x}_{item_j}$$

$$\mathbf{x}'_{user_i} = \oplus_{j \in \mathcal{N}(i)} \mathbf{x}_{user_j}$$

2.3.2. SAGEConv[1]

SAGEConv (GraphSAGE Convolutional Layer) updates the embeddings of items or users by incorporating information from both the target entity and its neighbors. Unlike SimpleConv, SAGEConv employs learnable weight matrices to compute the updates, enabling more flexibility and expressiveness in the model. Specifically, for each item or user, denoted as i , its updated embedding, \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} , is calculated as a linear combination of its own embedding and a weighted average of its neighbors' embeddings. The weights for this combination are determined by the weight matrices \mathbf{W}_1 and \mathbf{W}_2 . The first weight matrix \mathbf{W}_1 operates directly on the target entity's embedding, while the second weight matrix \mathbf{W}_2 operates on the mean of its neighbors' embeddings.

$$\begin{aligned}\mathbf{x}'_{item_i} &= \mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{W}_2 \cdot \text{mean}_{j \in \mathcal{N}(i)} \mathbf{x}_{item_j} \\ \mathbf{x}'_{user_i} &= \mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{W}_2 \cdot \text{mean}_{j \in \mathcal{N}(i)} \mathbf{x}_{user_j}\end{aligned}$$

2.3.3. GraphConv[2]

Similarly to SAGEConv, GraphConv (Graph Convolutional Layer) iteratively refines the embeddings of items or users by assimilating insights from both the focal entity and its surrounding neighbors within the graph. Additionally, GraphConv also integrates learnable weight matrices, \mathbf{W}_1 and \mathbf{W}_2 , into its computation. However, unlike SAGEConv each item or user represented by i , its updated embedding, \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} , is derived as a weighted sum of its own embedding and the aggregate of its neighbors' embeddings, instead of a weighted average.

$$\begin{aligned}\mathbf{x}'_{item_i} &= \mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{W}_2 \cdot \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{item_j} \\ \mathbf{x}'_{user_i} &= \mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{W}_2 \cdot \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{user_j}\end{aligned}$$

2.3.4. ResGatedGraphConv[3]

ResGatedGraphConv (Residual Gated Graph Convolutional Layer) extends the conventional graph convolution by incorporating residual connections and gating mechanisms. The updated embeddings of items or users, denoted as \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} , are computed through a combination of the original embeddings and the aggregated information from their neighboring entities within the graph. Specifically, for each item or user represented by i , the updated embedding is calculated as a weighted sum of its original embedding and the transformed embeddings of its neighbors, where the weighting factor $\eta_{i,j}$ is determined by a gating mechanism. This gating mechanism, governed by the sigmoid function σ , modulates the importance of neighboring embeddings based on their compatibility with the target entity's embedding, as encoded by the weight matrices \mathbf{W}_3 and \mathbf{W}_4 .

$$\begin{aligned}\mathbf{x}'_{item_i} &= \mathbf{W}_1 \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \eta_{i,j} \odot \mathbf{W}_2 \mathbf{x}_{item_j} \\ \eta_{i,j} &= \sigma(\mathbf{W}_3 \mathbf{x}_{item_i} + \mathbf{W}_4 \mathbf{x}_{item_j}) \\ \mathbf{x}'_{user_i} &= \mathbf{W}_1 \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \eta_{i,j} \odot \mathbf{W}_2 \mathbf{x}_{user_j} \\ \eta_{i,j} &= \sigma(\mathbf{W}_3 \mathbf{x}_{user_i} + \mathbf{W}_4 \mathbf{x}_{user_j})\end{aligned}$$

2.3.5. GATConv[4]

GATConv (Graph Attention Network Convolutional Layer) is designed to capture complex dependencies and attention mechanisms within the graph. The updated embeddings of items

or users, denoted as \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} respectively, are computed using attention mechanisms to dynamically weight the contributions of neighboring embeddings. Specifically, for each item or user represented by i , the updated embedding is a linear combination of the original embedding and the weighted sum of the embeddings of its neighbors. The attention coefficients $\alpha_{i,j}$ are computed using a shared self-attention mechanism across all neighbors, where the attention weights are determined by the attention mechanism applied to the pair of embeddings \mathbf{x}_{item_i} and \mathbf{x}_{item_j} (or \mathbf{x}_{user_i} and \mathbf{x}_{user_j}). The attention mechanism employs a LeakyReLU activation with a negative slope of 0.2 on a linear combination of the embeddings followed by a softmax function to ensure that the attention coefficients sum to one.

$$\begin{aligned}\mathbf{x}'_{item_i} &= \alpha_{i,i} \mathbf{W}_1 \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{item_j} \\ \alpha_{i,j} &= \frac{\exp(\text{LeakyReLU}(\mathbf{a}_1^\top \mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{a}_2^\top \mathbf{W}_2 \mathbf{x}_{item_j}))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\text{LeakyReLU}(\mathbf{a}_1^\top \mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{a}_2^\top \mathbf{W}_2 \mathbf{x}_{item_k}))} \\ \mathbf{x}'_{user_i} &= \alpha_{i,i} \mathbf{W}_1 \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{user_j} \\ \alpha_{i,j} &= \frac{\exp(\text{LeakyReLU}(\mathbf{a}_1^\top \mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{a}_2^\top \mathbf{W}_2 \mathbf{x}_{user_j}))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\text{LeakyReLU}(\mathbf{a}_1^\top \mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{a}_2^\top \mathbf{W}_2 \mathbf{x}_{user_k}))}\end{aligned}$$

2.3.6. GATv2Conv[5]

GATv2Conv (Graph Attention Network v2 Convolutional Layer) is an improvement of GATConv aimed at solving static attention coefficients. Similarly, GATv2Conv employs attention mechanisms to dynamically weigh the contributions of neighboring embeddings in a graph. The updated embeddings of items or users, denoted as \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} respectively, are computed using attention coefficients $\alpha_{i,j}$ determined by a self-attention mechanism. These attention coefficients control the influence of neighboring embeddings in the computation of updated embeddings. However, GATv2Conv solves the problem of static attention coefficients by applying attention mechanism directly on the concatenation of node embeddings before applying the LeakyReLU activation function.

$$\begin{aligned}\mathbf{x}'_{item_i} &= \alpha_{i,i} \mathbf{W}_1 \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{item_j} \\ \alpha_{i,j} &= \frac{\exp(\mathbf{a}^\top \text{LeakyReLU}(\mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{W}_2 \mathbf{x}_{item_j}))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\mathbf{a}^\top \text{LeakyReLU}(\mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{W}_2 \mathbf{x}_{item_k}))} \\ \mathbf{x}'_{user_i} &= \alpha_{i,i} \mathbf{W}_1 \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{user_j} \\ \alpha_{i,j} &= \frac{\exp(\mathbf{a}^\top \text{LeakyReLU}(\mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{W}_2 \mathbf{x}_{user_j}))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\mathbf{a}^\top \text{LeakyReLU}(\mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{W}_2 \mathbf{x}_{user_k}))}\end{aligned}$$

2.3.7. TransformerConv[6]

TransformerConv (Transformer Convolutional Layer) introduces a novel convolutional operation inspired by the Transformer[26] architecture. The updated embeddings of items or users, denoted as \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} , are computed as a weighted sum of the original embeddings of neighboring entities, where the attention weights $\alpha_{i,j}$ are determined through a self-attention mechanism. This mechanism computes the attention scores between the embeddings of the target entity and its neighbors using trainable weight matrices \mathbf{W}_3 and \mathbf{W}_4 , followed by a softmax operation to ensure a valid probability distribution. Additionally, a scaling factor \sqrt{d} is applied to stabilize the gradients during training, where d represents the dimensionality of the embeddings.

$$\mathbf{x}'_{item_i} = \mathbf{W}_1 \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{item_j}$$

$$\begin{aligned}
\alpha_{i,j} &= \text{softmax} \left(\frac{(\mathbf{W}_3 \mathbf{x}_{item_i})^\top (\mathbf{W}_4 \mathbf{x}_{item_j})}{\sqrt{d}} \right) \\
\mathbf{x}'_{user_i} &= \mathbf{W}_1 \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{user_j} \\
\alpha_{i,j} &= \text{softmax} \left(\frac{(\mathbf{W}_3 \mathbf{x}_{user_i})^\top (\mathbf{W}_4 \mathbf{x}_{user_j})}{\sqrt{d}} \right)
\end{aligned}$$

2.3.8. GINConv[7]

GINConv (Graph Isomorphism Network Convolutional Layer) introduces a convolutional operation that leverages neural networks. The updated embeddings of items or users, denoted as \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} respectively, are computed using a neural network NN_{Θ} with learnable parameters Θ . The operation involves a combination of the original embeddings of the target entity and its neighbors, scaled by a factor $(1 + \epsilon)$, where ϵ is a small constant.

$$\begin{aligned}
\mathbf{x}'_{item_i} &= NN_{\Theta} \left((1 + \epsilon) \cdot \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{item_j} \right) \\
\mathbf{x}'_{user_i} &= NN_{\Theta} \left((1 + \epsilon) \cdot \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{user_j} \right)
\end{aligned}$$

2.3.9. EdgeConv[8]

Similarly to GINConv, EdgeConv (Edge Convolutional Layer) introduces a convolutional operation that leverages neural networks. The updated embeddings of items or users, denoted as \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} respectively, are computed by aggregating information from neighboring entities within the graph. For each item or user represented by i , the updated embedding is obtained by summing over the embeddings of its neighbors, where each neighbor's embedding is transformed through a neural network NN_{Θ} with learnable parameters Θ . However, unlike GINConv the transformation involves concatenating the original embedding of the target entity with the difference between the neighbor's embedding and the target entity's embedding.

$$\begin{aligned}
\mathbf{x}'_{item_i} &= \sum_{j \in \mathcal{N}(i)} NN_{\Theta}(\mathbf{x}_{item_i} \parallel \mathbf{x}_{item_j} - \mathbf{x}_{item_i}) \\
\mathbf{x}'_{user_i} &= \sum_{j \in \mathcal{N}(i)} NN_{\Theta}(\mathbf{x}_{user_i} \parallel \mathbf{x}_{user_j} - \mathbf{x}_{user_i})
\end{aligned}$$

2.3.10. FeaStConv[9]

FeaStConv (Feature-Steered Convolutional Layer) is another convolutional operation that attempts to employ an attention mechanism. The updated embeddings of items or users, denoted as \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} respectively, are computed by aggregating information from neighboring entities within the graph. Specifically, for each item or user represented by i , the updated embedding is obtained by averaging over the embeddings of its neighbors, weighted by a feature-steering mechanism $q_h(\mathbf{x}_{item_i}, \mathbf{x}_{item_j})$ or $q_h(\mathbf{x}_{user_i}, \mathbf{x}_{user_j})$ for items or users, respectively, where h denotes an attention head. This mechanism, parameterized by learnable vectors \mathbf{u}_h and biases c_h , steers the convolutional process towards relevant features by computing attention scores based on the feature differences between the target entity and its neighbors. The weighted sum is then multiplied by learnable weight matrices \mathbf{W}_h to produce the final updated embedding.

$$\begin{aligned}
\mathbf{x}'_{item_i} &= \frac{1}{|\mathcal{N}(i)|} \sum_{j \in \mathcal{N}(i)} \sum_{h=1}^H q_h(\mathbf{x}_{item_i}, \mathbf{x}_{item_j}) \mathbf{W}_h \mathbf{x}_{item_j} \\
q_h(\mathbf{x}_{item_i}, \mathbf{x}_{item_j}) &= \text{softmax}_j(\mathbf{u}_h^\top (\mathbf{x}_{item_j} - \mathbf{x}_{item_i}) + c_h)
\end{aligned}$$

$$\mathbf{x}'_{user_i} = \frac{1}{|\mathcal{N}(i)|} \sum_{j \in \mathcal{N}(i)} \sum_{h=1}^H q_h(\mathbf{x}_{user_i}, \mathbf{x}_{user_j}) \mathbf{W}_h \mathbf{x}_{user_j}$$

$$q_h(\mathbf{x}_{user_i}, \mathbf{x}_{user_j}) = \text{softmax}_j(\mathbf{u}_h^\top (\mathbf{x}_{user_j} - \mathbf{x}_{user_i}) + c_h)$$

2.3.11. LEConv[10]

LEConv (Linear Embedding Convolutional Layer) is a convolutional operation that focuses on capturing linear relationships between entities within the graph. The updated embeddings of items or users, denoted as \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} respectively, are computed by combining the original embeddings with the differences between the transformed embeddings of the target entity and its neighbors. For each item or user represented by i , the updated embedding is obtained by multiplying the original embedding with \mathbf{W}_1 and summing over the differences between the transformed embeddings of the target entity and its neighbors, weighted by \mathbf{W}_2 and \mathbf{W}_3 .

$$\mathbf{x}'_{item_i} = \mathbf{W}_1 \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} (\mathbf{W}_2 \mathbf{x}_{item_i} - \mathbf{W}_3 \mathbf{x}_{item_j})$$

$$\mathbf{x}'_{user_i} = \mathbf{W}_1 \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} (\mathbf{W}_2 \mathbf{x}_{user_i} - \mathbf{W}_3 \mathbf{x}_{user_j})$$

2.3.12. GENConv[11]

GENConv (Graph Edge Network Convolutional Layer) is a convolutional operation that leverages multi-layered perceptrons (MLPs). The updated embeddings of items or users, denoted as \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} respectively, are computed by applying a MLP to the sum of the original embedding of the target entity and the ReLU-activated sum of the embeddings of its neighbors, where ϵ is a small constant for numerical stability.

$$\mathbf{x}'_{item_i} = \text{MLP} \left(\mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \text{ReLU}(\mathbf{x}_{item_j}) + \epsilon \right)$$

$$\mathbf{x}'_{user_i} = \text{MLP} \left(\mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \text{ReLU}(\mathbf{x}_{user_j}) + \epsilon \right)$$

2.3.13. WLConvContinuous[12]

WLConvContinuous (Weisfeiler-Lehman Convolutional Layer with Continuous Weights) is a convolutional operation that incorporates Weisfeiler-Lehman graph isomorphism test and continuous weighting mechanisms. The updated embeddings of items or users, denoted as \mathbf{x}'_{item_i} or \mathbf{x}'_{user_i} respectively, are computed by averaging the original embeddings of the target entity and its neighbors, weighted by the inverse of their degrees. Specifically, for each item or user represented by i , the updated embedding is obtained by taking half of the original embedding and half of the average of the embeddings of its neighbors, weighted by $\frac{1}{\deg(i)}$. This weighting scheme ensures that entities with higher degrees contribute proportionally less to the final embedding, mitigating the influence of highly connected nodes in the graph.

$$\mathbf{x}'_{item_i} = \frac{1}{2} \left(\mathbf{x}_{item_i} + \frac{1}{\deg(i)} \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{item_j} \right)$$

$$\mathbf{x}'_{user_i} = \frac{1}{2} \left(\mathbf{x}_{user_i} + \frac{1}{\deg(i)} \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{user_j} \right)$$

2.4. Dataset

The MovieLens100k dataset is a classic benchmark dataset widely used in recommendation system research. It consists of 100,000 timestamped ratings (from 1 to 5) given by users to approximately 1,000 movies. The dataset is composed of 4 main features, these are: user id's,

movie id’s, ratings and timestamps. Therefore, the corresponding rating, given by each user to a particular movie, will serve as our edge label, or target \mathbf{t} , which we will try to approximate using the previously mentioned bipartite GNN operators.

For the purposes of this benchmark, we preprocessed ratings to include only those equal to or greater than 4, ensuring a focus on high-quality user preferences and emphasizing positive interactions with movies. A description of the resulting dataset after being preprocessed can be seen in table 1.

Dataset	Users	Items	Interactions
MovieLens100k[27]	942	1447	55,375

Table 1: Preprocessed MovieLens100k Description.

3. Results

Both the figures and appendix tables C present a comprehensive overview of the performance metrics of each bipartite GNN model applied to MovieLens100k dataset across different evaluation metrics, namely Recall@ k , Precision@ k , MAP@ k , F1@ k , and nDCG@ k . Each figure and table focuses on a specific metric, providing results for different values of k , representing the top- k recommendations. Overall, by first glance of the figures, the majority of the bipartite GNN models seem to have performed relatively the same, however there are three clear models that seem to have underperformed, these are: SimpleConv, WLConvContinuous and EdgeConv. The performance of the first two of the previously mentioned models was to be expected since they are comprised of aggregating input embeddings. Nevertheless, the underperformance of EdgeConv was not expected since the convolutional operator leverages neural networks.

In figure 2 and table 4, the results for Recall@ k are displayed. Notably, the best performing models were SageConv, ResGatedGraphConv, TransformerConv and GINConv. However, from previously mentioned models TransformerConv had the most success when it came to achieving top scores for Recall@ k , demonstrating its ability to dynamically recommend top- k items from a list relevant items greater than 70.

In figure 3 and table 5, the results for Precision@ k are displayed. Notably, the best performing models were SageConv, GraphConv, ResGatedGraphConv, GINConv and FeaStConv. However, from the previously mentioned models SAGEConv had the most success when it came to achieving top scores for Precision@ k , demonstrating its ability to dynamically recommend top- k items from a limited list relevant items roughly greater than 50.

In figure 4 and table 6, the results for MAP@ k are displayed. Notably, GINConv and LEConv. However, from the previously mentioned models GINConv significantly stood over the rest when it came to achieving top scores for MAP@ k , demonstrating its ability to dynamically rank relevant items across all users or queries.

In figure 5 and table 7, the results for F1@ k are displayed. Notably, SageConv, GraphConv, ResGatedGraphConv, TransformerConv and GINConv. However, from the previously mentioned models SageConv significantly stood over the rest when it came to achieving top scores for F1@ k , demonstrating a weighted ability to both dynamically provide and rank relevant items across all users or queries.

In figure 6 and table 8, the results for nDCG@ k are displayed. Notably, GINConv and LEConv. However, from the previously mentioned models GINConv significantly stood over the rest when it came to achieving top scores for nDCG@ k , demonstrating its ability to position dynamically relevant items across all users or queries.

Lastly, apart from the visually underperforming models mentioned at the beginning of this section, other models such as GATConv, GATv2Conv and GENConv were also unable to get a single top score for each metric at any given k .

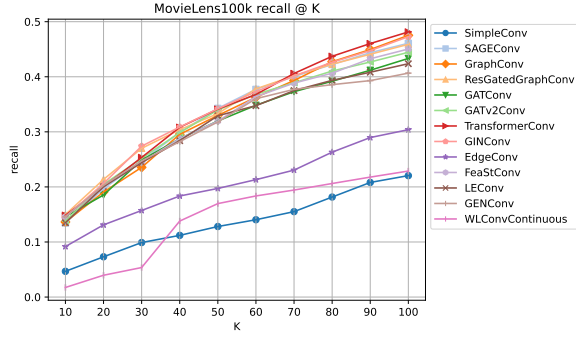


Figure 2: MovieLens100k Results for Recall@ k .




Figure 3: MovieLens100k Results for Precision@ k .




Figure 4: MovieLens100k Results for MAP@ k .




Figure 5: MovieLens100k Results for F1@ k .




Figure 6: MovieLens100k Results for nDCG@ k .

4. Discussion

In table 9, an aggregation of the top-performing models across all metrics is presented. This table provides a summary of the most successful models based on their cumulative performance across the evaluated metrics. Noticeably, from the total column, we can observe that the two best models were GINConv with 22 top scores and SAGEConv with 11 top scores. This superior performance in our dynamic recommendation system benchmark can be attributed to their ability to effectively leverage graph structure, adaptively combine information from the target entity and its neighbors, and introduce flexibility and expressiveness through learnable parameters. These attributes enable GINConv and SAGEConv to capture intricate relationships and dependencies within the recommendation system, ultimately leading to improved recommendation performance.

For future work, we propose several avenues to further enhance dynamic recommendation systems. Firstly, exploring deeper architectures of Graph Neural Networks (GNNs) beyond GINConv and SAGEConv could unlock additional improvements in recommendation accuracy and personalization. Secondly, integrating Bipartite LSTMs into the recommendation framework offers a promising approach to modeling sequential interactions in bipartite graphs, potentially capturing temporal dynamics and long-term user preferences more effectively. Additionally, expanding the scope of datasets used for benchmarking could provide a more comprehensive evaluation of recommendation algorithms across diverse domains and scenarios. Lastly, leveraging more advanced computing resources such as AWS EC2 instances could facilitate extensive experimentation and model training, enabling us to tackle larger-scale recommendation problems and explore more complex architectures effectively.

5. Conclusion

In this research, we embarked on a meticulous exploration of Graph Neural Networks (GNNs) applied to dynamic recommendation systems within the context of undirected bipartite graphs. Our primary objective was to establish a standardized benchmarking framework to systematically evaluate the efficacy of various GNN operators tailored for this domain. To achieve this, we meticulously curated a comprehensive set of GNN models ranging from traditional techniques to state-of-the-art architectures, leveraging the PyTorch Geometric library for implementation.

Our contributions were threefold. Firstly, we developed a flexible and extensible benchmarking framework that encompasses data preprocessing, model training, evaluation protocols, and result analysis. This framework enables fair comparison and evaluation of GNN-based recommendation models under different experimental settings. Secondly, we rigorously evaluated the performance of each GNN operator on the MovieLens100k dataset, shedding light on their strengths and weaknesses in capturing intricate user-item interactions. Lastly, through insightful analysis of experimental results, we provided practical recommendations for selecting and configuring GNN models tailored to specific application domains, aiming to empower researchers and practitioners in the field.

By offering this comprehensive benchmarking study along with an open-source implementation in PyTorch Geometric, we aim to foster transparency, reproducibility, and innovation in the development of GNN-based dynamic recommendation systems. We believe that our work serves as a valuable resource for the community, facilitating advancements in recommendation engine technology and ultimately enhancing user experiences across various domains.

Appendix A Summary of Models

Model	Equations
SimpleConv	$\mathbf{x}'_{item_i} = \bigoplus_{j \in \mathcal{N}(i)} \mathbf{x}_{item_j}$ $\mathbf{x}'_{user_i} = \bigoplus_{j \in \mathcal{N}(i)} \mathbf{x}_{user_j}$
SAGEConv	$\mathbf{x}'_{item_i} = \mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{W}_2 \cdot \text{mean}_{j \in \mathcal{N}(i)} \mathbf{x}_{item_j}$ $\mathbf{x}'_{user_i} = \mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{W}_2 \cdot \text{mean}_{j \in \mathcal{N}(i)} \mathbf{x}_{user_j}$
GraphConv	$\mathbf{x}'_{item_i} = \mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{W}_2 \cdot \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{item_j}$ $\mathbf{x}'_{user_i} = \mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{W}_2 \cdot \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{user_j}$
ResGatedGraphConv	$\mathbf{x}'_{item_i} = \mathbf{W}_1 \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \eta_{i,j} \odot \mathbf{W}_2 \mathbf{x}_{item_j}$ $\eta_{i,j} = \sigma(\mathbf{W}_3 \mathbf{x}_{item_i} + \mathbf{W}_4 \mathbf{x}_{item_j})$ $\mathbf{x}'_{user_i} = \mathbf{W}_1 \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \eta_{i,j} \odot \mathbf{W}_2 \mathbf{x}_{user_j}$ $\eta_{i,j} = \sigma(\mathbf{W}_3 \mathbf{x}_{user_i} + \mathbf{W}_4 \mathbf{x}_{user_j})$
GATConv	$\mathbf{x}'_{item_i} = \alpha_{i,i} \mathbf{W}_1 \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{item_j}$ $\alpha_{i,j} = \frac{\exp(\text{LeakyReLU}(\mathbf{a}_1^\top \mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{a}_2^\top \mathbf{W}_2 \mathbf{x}_{item_j}))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\text{LeakyReLU}(\mathbf{a}_1^\top \mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{a}_2^\top \mathbf{W}_2 \mathbf{x}_{item_k}))}$ $\mathbf{x}'_{user_i} = \alpha_{i,i} \mathbf{W}_1 \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{user_j}$ $\alpha_{i,j} = \frac{\exp(\text{LeakyReLU}(\mathbf{a}_1^\top \mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{a}_2^\top \mathbf{W}_2 \mathbf{x}_{user_j}))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\text{LeakyReLU}(\mathbf{a}_1^\top \mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{a}_2^\top \mathbf{W}_2 \mathbf{x}_{user_k}))}$
GATv2Conv	$\mathbf{x}'_{item_i} = \alpha_{i,i} \mathbf{W}_1 \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{item_j}$ $\alpha_{i,j} = \frac{\exp(\mathbf{a}^\top \text{LeakyReLU}(\mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{W}_2 \mathbf{x}_{item_j}))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\mathbf{a}^\top \text{LeakyReLU}(\mathbf{W}_1 \mathbf{x}_{item_i} + \mathbf{W}_2 \mathbf{x}_{item_k}))}$ $\mathbf{x}'_{user_i} = \alpha_{i,i} \mathbf{W}_1 \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{user_j}$ $\alpha_{i,j} = \frac{\exp(\mathbf{a}^\top \text{LeakyReLU}(\mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{W}_2 \mathbf{x}_{user_j}))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\mathbf{a}^\top \text{LeakyReLU}(\mathbf{W}_1 \mathbf{x}_{user_i} + \mathbf{W}_2 \mathbf{x}_{user_k}))}$
TransformerConv	$\mathbf{x}'_{item_i} = \mathbf{W}_1 \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{item_j}$ $\alpha_{i,j} = \text{softmax} \left(\frac{(\mathbf{W}_3 \mathbf{x}_{item_i})^\top (\mathbf{W}_4 \mathbf{x}_{item_j})}{\sqrt{d}} \right)$ $\mathbf{x}'_{user_i} = \mathbf{W}_1 \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \mathbf{W}_2 \mathbf{x}_{user_j}$

Model	Equations
	$\alpha_{i,j} = \text{softmax} \left(\frac{(\mathbf{W}_3 \mathbf{x}_{user_i})^\top (\mathbf{W}_4 \mathbf{x}_{user_j})}{\sqrt{d}} \right)$
GINConv	$\mathbf{x}'_{item_i} = NN_{\Theta} \left((1 + \epsilon) \cdot \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{item_j} \right)$ $\mathbf{x}'_{user_i} = NN_{\Theta} \left((1 + \epsilon) \cdot \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{user_j} \right)$
EdgeConv	$\mathbf{x}'_{item_i} = \sum_{j \in \mathcal{N}(i)} NN_{\Theta}(\mathbf{x}_{item_i} \parallel \mathbf{x}_{item_j} - \mathbf{x}_{item_i})$ $\mathbf{x}'_{user_i} = \sum_{j \in \mathcal{N}(i)} NN_{\Theta}(\mathbf{x}_{user_i} \parallel \mathbf{x}_{user_j} - \mathbf{x}_{user_i})$
FeaStConv	$\mathbf{x}'_{item_i} = \frac{1}{ \mathcal{N}(i) } \sum_{j \in \mathcal{N}(i)} \sum_{h=1}^H q_h(\mathbf{x}_{item_i}, \mathbf{x}_{item_j}) \mathbf{W}_h \mathbf{x}_{item_j}$ $q_h(\mathbf{x}_{item_i}, \mathbf{x}_{item_j}) = \text{softmax}_j(\mathbf{u}_h^\top (\mathbf{x}_{item_j} - \mathbf{x}_{item_i}) + c_h)$ $\mathbf{x}'_{user_i} = \frac{1}{ \mathcal{N}(i) } \sum_{j \in \mathcal{N}(i)} \sum_{h=1}^H q_h(\mathbf{x}_{user_i}, \mathbf{x}_{user_j}) \mathbf{W}_h \mathbf{x}_{user_j}$ $q_h(\mathbf{x}_{user_i}, \mathbf{x}_{user_j}) = \text{softmax}_j(\mathbf{u}_h^\top (\mathbf{x}_{user_j} - \mathbf{x}_{user_i}) + c_h)$
LEConv	$\mathbf{x}'_{item_i} = \mathbf{W}_1 \mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} (\mathbf{W}_2 \mathbf{x}_{item_i} - \mathbf{W}_3 \mathbf{x}_{item_j})$ $\mathbf{x}'_{user_i} = \mathbf{W}_1 \mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} (\mathbf{W}_2 \mathbf{x}_{user_i} - \mathbf{W}_3 \mathbf{x}_{user_j})$
GENConv	$\mathbf{x}'_{item_i} = \text{MLP} \left(\mathbf{x}_{item_i} + \sum_{j \in \mathcal{N}(i)} \text{ReLU}(\mathbf{x}_{item_j}) + \epsilon \right)$ $\mathbf{x}'_{user_i} = \text{MLP} \left(\mathbf{x}_{user_i} + \sum_{j \in \mathcal{N}(i)} \text{ReLU}(\mathbf{x}_{user_j}) + \epsilon \right)$
WLConvContinuous	$\mathbf{x}'_{item_i} = \frac{1}{2} \left(\mathbf{x}_{item_i} + \frac{1}{\deg(i)} \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{item_j} \right)$ $\mathbf{x}'_{user_i} = \frac{1}{2} \left(\mathbf{x}_{user_i} + \frac{1}{\deg(i)} \sum_{j \in \mathcal{N}(i)} \mathbf{x}_{user_j} \right)$

Table 2: Summary of Bipartite GNN Models.

Appendix B Summary of Hyperparameters

Model	Bipartite GNN Layers	Linear Layers	Neural Networks	Heads
SimpleConv	1	1	-	-
SAGEConv	2	2	-	-
GraphConv	2	2	-	-
ResGatedGraphConv	2	2	-	-
GATConv	3	3	-	8
GATv2Conv	3	3	-	8
TransformerConv	3	-	-	8
GINConv	1	1	1MLP-3L	-
EdgeConv	1	1	1MLP-2L	-
FeaStConv	2	3	-	2
LEConv	3	3	-	-
GENConv	3	3	-	-
WLConvContinuous	3	3	-	-

Table 3: Summary of Bipartite GNN Model Hyperparameters.

Appendix C Summary of Results

Model	K									
	10	20	30	40	50	60	70	80	90	100
SimpleConv	0.0467	0.0731	0.0990	0.1120	0.1281	0.1404	0.1551	0.1816	0.2081	0.2205
SAGEConv	0.1346	0.1965	0.2471	0.2989	0.3426	0.3771	0.4004	0.4223	0.4437	0.4608
GraphConv	0.1364	0.1907	0.2355	0.2955	0.3298	0.3628	0.3932	0.4274	0.4494	0.4752
ResGatedGraphConv	0.1498	0.2136	0.2701	0.3040	0.3330	0.3780	0.4018	0.4228	0.4409	0.4589
GATConv	0.1461	0.1853	0.2488	0.2876	0.3197	0.3482	0.3729	0.3919	0.4118	0.4333
GATv2Conv	0.1381	0.2045	0.2518	0.2987	0.3396	0.3682	0.3890	0.4104	0.4268	0.4444
TransformerConv	0.1485	0.2003	0.2535	0.3083	0.3412	0.3678	0.4060	0.4372	0.4603	0.4814
GINConv	0.1458	0.2062	0.2740	0.3088	0.3419	0.3729	0.4014	0.4264	0.4475	0.4728
EdgeConv	0.0917	0.1311	0.1572	0.1835	0.1972	0.2132	0.2305	0.2631	0.2896	0.3039
FeaStConv	0.1458	0.2007	0.2446	0.2878	0.3194	0.3639	0.3885	0.4065	0.4327	0.4504
LEConv	0.1350	0.2004	0.2441	0.2840	0.3292	0.3475	0.3749	0.3935	0.4082	0.4237
GENConv	0.1457	0.2028	0.2473	0.2822	0.3180	0.3602	0.3771	0.3857	0.3930	0.4066
WLConvContinuous	0.0174	0.0396	0.0536	0.1380	0.1697	0.1836	0.1943	0.2062	0.2178	0.2288

Table 4: MovieLens100k Results for Recall@ k .

Model	K									
	10	20	30	40	50	60	70	80	90	100
SimpleConv	0.1107	0.1236	0.1100	0.0936	0.0858	0.0793	0.0744	0.0737	0.0748	0.0748
SAGEConv	0.2264	0.1846	0.1723	0.1663	0.1542	0.1435	0.1357	0.1302	0.1248	0.1210
GraphConv	0.2279	0.1777	0.1555	0.1554	0.1470	0.1427	0.1353	0.1294	0.1247	0.1213
ResGatedGraphConv	0.2404	0.1904	0.1783	0.1600	0.1464	0.1359	0.1278	0.1208	0.1182	0.1155
GATConv	0.2425	0.1689	0.1625	0.1512	0.1419	0.1352	0.1288	0.1209	0.1179	0.1135
GATv2Conv	0.2261	0.1946	0.1751	0.1530	0.1439	0.1371	0.1303	0.1248	0.1178	0.1139
TransformerConv	0.2457	0.1838	0.1664	0.1578	0.1447	0.1360	0.1286	0.1257	0.1223	0.1178
GINConv	0.2668	0.1966	0.1780	0.1562	0.1469	0.1394	0.1317	0.1268	0.1238	0.1198
EdgeConv	0.1361	0.1386	0.1240	0.1074	0.0976	0.0895	0.0839	0.0825	0.0829	0.0824
FeaStConv	0.2446	0.1848	0.1718	0.1669	0.1505	0.1437	0.1342	0.1257	0.1229	0.1177
LEConv	0.2286	0.1880	0.1611	0.1527	0.1427	0.1298	0.1267	0.1199	0.1094	0.1031
GENConv	0.2421	0.1820	0.1590	0.1473	0.1350	0.1264	0.1145	0.1068	0.0987	0.0939
WLConvContinuous	0.0364	0.0463	0.0475	0.0797	0.0791	0.0742	0.0679	0.0646	0.0664	0.0654

Table 5: MovieLens100k Results for Precision@ k .

Model	K									
	10	20	30	40	50	60	70	80	90	100
SimpleConv	0.0382	0.0417	0.0395	0.0351	0.0336	0.0326	0.0317	0.0318	0.0327	0.0334
SAGEConv	0.1527	0.1324	0.1315	0.1410	0.1465	0.1489	0.1526	0.1539	0.1567	0.1605
GraphConv	0.1528	0.1257	0.1240	0.1392	0.1424	0.1452	0.1495	0.1527	0.1564	0.1581
ResGatedGraphConv	0.1698	0.1390	0.1395	0.1398	0.1410	0.1517	0.1546	0.1566	0.1592	0.1613
GATConv	0.1574	0.1210	0.1277	0.1304	0.1316	0.1355	0.1392	0.1425	0.1480	0.1508
GATv2Conv	0.1581	0.1379	0.1345	0.1397	0.1443	0.1467	0.1486	0.1507	0.1530	0.1552
TransformerConv	0.1568	0.1286	0.1291	0.1405	0.1421	0.1442	0.1482	0.1527	0.1561	0.1580
GINConv	0.1778	0.1358	0.1502	0.1496	0.1525	0.1550	0.1578	0.1602	0.1624	0.1645
EdgeConv	0.0791	0.0809	0.0803	0.0765	0.0738	0.0745	0.0728	0.0733	0.0742	0.0751
FeaStConv	0.1622	0.1346	0.1367	0.1412	0.1448	0.1530	0.1545	0.1562	0.1598	0.1589
LEConv	0.1803	0.1483	0.1423	0.1436	0.1468	0.1479	0.1496	0.1518	0.1518	0.1531
GENConv	0.1738	0.1297	0.1196	0.1173	0.1037	0.1046	0.0998	0.0993	0.0964	0.0982
WLConvContinuous	0.0068	0.0084	0.0098	0.0392	0.0618	0.0727	0.0762	0.0804	0.0842	0.0858

Table 6: MovieLens100k Results for MAP@ k .

Model	K									
	10	20	30	40	50	60	70	80	90	100
SimpleConv	0.0490	0.0693	0.0752	0.0751	0.0754	0.0751	0.0758	0.0807	0.0865	0.0888
SAGEConv	0.1295	0.1465	0.1584	0.1676	0.1691	0.1672	0.1645	0.1623	0.1595	0.1581
GraphConv	0.1309	0.1413	0.1472	0.1617	0.1624	0.1648	0.1631	0.1606	0.1589	0.1585
ResGatedGraphConv	0.1416	0.1554	0.1671	0.1645	0.1617	0.1595	0.1568	0.1531	0.1524	0.1518
GATConv	0.1414	0.1383	0.1547	0.1572	0.1571	0.1548	0.1543	0.1508	0.1506	0.1483
GATv2Conv	0.1313	0.1543	0.1612	0.1592	0.1616	0.1602	0.1578	0.1556	0.1510	0.1493
TransformerConv	0.1432	0.1481	0.1553	0.1623	0.1606	0.1584	0.1574	0.1588	0.1580	0.1560
GINConv	0.1429	0.1550	0.1683	0.1650	0.1644	0.1635	0.1609	0.1600	0.1592	0.1574
EdgeConv	0.0738	0.0901	0.0938	0.0939	0.0920	0.0896	0.0895	0.0937	0.0985	0.1002
FeaStConv	0.1413	0.1489	0.1571	0.1653	0.1617	0.1646	0.1603	0.1557	0.1559	0.1531
LEConv	0.1310	0.1518	0.1541	0.1586	0.1592	0.1523	0.1529	0.1493	0.1406	0.1362
GENConv	0.1392	0.1458	0.1497	0.1505	0.1483	0.1472	0.1406	0.1347	0.1274	0.1241
WLConvContinuous	0.0127	0.0317	0.0375	0.0781	0.0835	0.0824	0.0790	0.0780	0.0800	0.0805

Table 7: MovieLens100k Results for F1@ k .

Model	K									
	10	20	30	40	50	60	70	80	90	100
SimpleConv	0.1004	0.1190	0.1210	0.1173	0.1185	0.1198	0.1226	0.1313	0.1413	0.1470
SAGEConv	0.2634	0.2604	0.2733	0.2976	0.3135	0.3239	0.3328	0.3410	0.3489	0.3580
GraphConv	0.2606	0.2484	0.2556	0.2896	0.3027	0.3165	0.3268	0.3392	0.3477	0.3585
ResGatedGraphConv	0.2832	0.2691	0.2837	0.2911	0.2983	0.3203	0.3291	0.3366	0.3452	0.3541
GATConv	0.2699	0.2396	0.2633	0.2759	0.2862	0.3008	0.3128	0.3209	0.3327	0.3412
GATv2Conv	0.2689	0.2677	0.2759	0.2922	0.3057	0.3172	0.3253	0.3331	0.3382	0.3470
TransformerConv	0.2724	0.2522	0.2672	0.2948	0.3060	0.3156	0.3276	0.3406	0.3507	0.3589
GINConv	0.2966	0.2656	0.2980	0.3023	0.3140	0.3243	0.3342	0.3437	0.3534	0.3618
EdgeConv	0.1560	0.1745	0.1790	0.1805	0.1799	0.1830	0.1857	0.1955	0.2058	0.2122
FeaStConv	0.2773	0.2613	0.2791	0.2962	0.3049	0.3234	0.3310	0.3371	0.3474	0.3526
LEConv	0.2881	0.2748	0.2783	0.2901	0.3039	0.3093	0.3190	0.3270	0.3277	0.3343
GENConv	0.2910	0.2559	0.2572	0.2635	0.2576	0.2681	0.2660	0.2677	0.2646	0.2706
WLConvContinuous	0.0262	0.0426	0.0502	0.1258	0.1656	0.1824	0.1886	0.1958	0.2065	0.2117

Table 8: MovieLens100k Results for nDCG@ k .

Model	Metrics					Total
	Recall	Precision	MAP	F1	nDCG	
SimpleConv	-	-	-	-	-	-
SAGEConv	1	4	-	6	-	11
GraphConv	-	1	-	1	-	2
ResGatedGraphConv	3	1	-	1	-	5
GATConv	-	-	-	-	-	-
GATv2Conv	-	-	-	-	-	-
TransformerConv	4	-	-	1	-	5
GINConv	2	2	8	1	9	22
EdgeConv	-	-	-	-	-	-
FeaStConv	-	2	-	-	-	2
LEConv	-	-	2	-	1	3
GENConv	-	-	-	-	-	-
WLConvContinuous	-	-	-	-	-	-

Table 9: MovieLens100k Aggregated Top Results.

References

- [1] William L. Hamilton, Rex Ying, and Jure Leskovec. Inductive representation learning on large graphs, 2018.
- [2] Christopher Morris, Martin Ritzert, Matthias Fey, William L. Hamilton, Jan Eric Lenssen, Gaurav Rattan, and Martin Grohe. Weisfeiler and leman go neural: Higher-order graph neural networks, 2021.
- [3] Xavier Bresson and Thomas Laurent. Residual gated graph convnets, 2018.
- [4] Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. Graph attention networks, 2018.
- [5] Shaked Brody, Uri Alon, and Eran Yahav. How attentive are graph attention networks?, 2022.
- [6] Yunsheng Shi, Zhengjie Huang, Shikun Feng, Hui Zhong, Wenjin Wang, and Yu Sun. Masked label prediction: Unified message passing model for semi-supervised classification, 2021.
- [7] Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. How powerful are graph neural networks?, 2019.
- [8] Yue Wang, Yongbin Sun, Ziwei Liu, Sanjay E. Sarma, Michael M. Bronstein, and Justin M. Solomon. Dynamic graph cnn for learning on point clouds, 2019.
- [9] Nitika Verma, Edmond Boyer, and Jakob Verbeek. Feastnet: Feature-steered graph convolutions for 3d shape analysis, 2018.
- [10] Ekagra Ranjan, Soumya Sanyal, and Partha Pratim Talukdar. Asap: Adaptive structure aware pooling for learning hierarchical graph representations, 2020.
- [11] Guohao Li, Chenxin Xiong, Ali Thabet, and Bernard Ghanem. Deepergcnn: All you need to train deeper gcns, 2020.
- [12] Matteo Togninalli, Elisabetta Ghisu, Felipe Llinares-López, Bastian Rieck, and Karsten Borgwardt. Wasserstein weisfeiler-lehman graph kernels, 2019.
- [13] Hervé Jégou, Matthijs Douze, and Cordelia Schmid. Product quantization for nearest neighbor search. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 33(1):117–128, 2011.
- [14] Shicong Liu, Hongtao Lu, and Junru Shao. Improved residual vector quantization for high-dimensional approximate nearest neighbor search, 2015.
- [15] Wei Dong, Moses Charikar, and Kai Li. Efficient k-nearest neighbor graph construction for generic similarity measures. In *Proceedings of the 20th International Conference on World Wide Web*, pages 577–586. ACM, 2011.
- [16] Matthijs Douze, Alexandr Guzhva, Chengqi Deng, Jeff Johnson, Gergely Szilvasy, Pierre-Emmanuel Mazaré, Maria Lomeli, Lucas Hosseini, and Hervé Jégou. The faiss library, 2024.
- [17] Weihua Hu, Bowen Liu, Joseph Gomes, Marinka Zitnik, Percy Liang, Vijay Pande, and Jure Leskovec. Strategies for pre-training graph neural networks, 2020.
- [18] Federico Monti, Davide Boscaini, Jonathan Masci, Emanuele Rodolà, Jan Svoboda, and Michael M. Bronstein. Geometric deep learning on graphs and manifolds using mixture model cnns, 2016.

- [19] Matthias Fey, Jan Eric Lenssen, Frank Weichert, and Heinrich Müller. Splinecnn: Fast geometric deep learning with continuous b-spline kernels, 2018.
- [20] Justin Gilmer, Samuel S. Schoenholz, Patrick F. Riley, Oriol Vinyals, and George E. Dahl. Neural message passing for quantum chemistry, 2017.
- [21] Tian Xie and Jeffrey C. Grossman. Crystal graph convolutional neural networks for an accurate and interpretable prediction of material properties. *Phys. Rev. Lett.*, 120:145301, Apr 2018.
- [22] Jiaxuan You, Rex Ying, and Jure Leskovec. Design space for graph neural networks, 2021.
- [23] David Duvenaud, Dougal Maclaurin, Jorge Aguilera-Iparraguirre, Rafael Gómez-Bombarelli, Timothy Hirzel, Alán Aspuru-Guzik, and Ryan P. Adams. Convolutional networks on graphs for learning molecular fingerprints, 2015.
- [24] Deyu Bo, Xiao Wang, Chuan Shi, and Huawei Shen. Beyond low-frequency information in graph convolutional networks, 2021.
- [25] David Buterez, Jon Paul Janet, Steven J. Kiddle, Dino Oglic, and Pietro Liò. Graph neural networks with adaptive readouts, 2022.
- [26] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need, 2023.
- [27] F. Maxwell Harper and Joseph A. Konstan. The movielens datasets: History and context. *ACM Transactions on Interactive Intelligent Systems (TiiS)*, 5(4):19:1–19:19, 2015.